



Architecting and Operating OpenShift Clusters: OpenShift for Infrastructure and O..



PREV

9. Day-2 Opera



AA



NEXT



PI Mode Ba...

© William Caban 2019

W. Caban, *Architecting and Operating OpenShift Clusters*

[https://doi.org/10.1007/978-1-4842-4985-7\\_10](https://doi.org/10.1007/978-1-4842-4985-7_10)

# 10. Advanced Network Operations

William Caban

<sup>1</sup>

(1) Columbia, MD, USA



The OpenShift platform provides defaults optimized for Cloud-native workloads. These have been covered throughout this book. As with many successful Open Source project, Kubernetes is being used in setups for which it was never designed. With the adoption of Kubernetes outside the web-based application, there is the need to support specialized hardware acceleration (i.e., GPUs, FPGAs, etc.), multiple NICs per Container, and much more. This chapter focuses on advanced networking features or capabilities for increasing network performance and for the onboarding of applications or microservices using nontraditional web protocols into *OpenShift*.

---

## Network Optimizations

*OpenShift SDN* uses *OpenvSwitch*, *VXLAN* tunnels, *OpenFlow* rules, and *iptables* or *firewalld* rules. Some possible optimizations to this overlay network are based on best practices for fine-tuning a system in a high-performance environment.

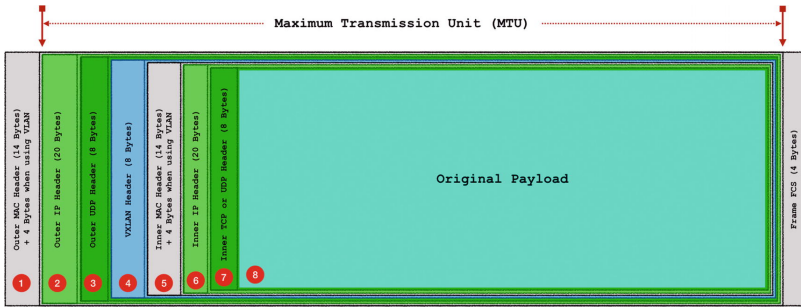
## JUMBO FRAMES AND VXLAN ACCELERATION

The standard Ethernet Maximum Transmission Unit (MTU) is 1500 Bytes. A regular IP UDP packet will consume 20 Bytes for the IP header (see #2 in Figure [10-1](#)) and 8 Bytes for the UDP header (see #3 in Figure [10-1](#)), and the remaining 1472 Bytes are available for payload (see #4 to #8 in Figure [10-1](#)).

**NOTE** The outer Ethernet header (14 Bytes) (see #1 in Figure [10-1](#)) is not counted as part of the MTU.

In SDN networks using the VXLAN protocol, the whole Ethernet frame of traffic from *Pods* in one *Node* destined to *Pods* in another *Node* is encapsulated as IP UDP packets and forwarded to the *Node* running the destination *Pods*. For this, the VXLAN header (see #4 in Figure [10-1](#)) is added to the original Layer 2 Ethernet frame (see #5 in Figure [10-1](#)), minus its FCS, and all this content becomes the payload of the outer IP UDP packets (see #2 and #3 in Figure [10-1](#)) and is sent to the remote *Node*.





**Figure 10-1** Format of a VXLAN packet

From the diagram in Figure [10-1](#), the effective MTU for payload (see #8) is 1422 Bytes. When working with environments with large streams of data to transfer among Pods on different Nodes, those streams of data need to be broken into very small chunks of 1422 Bytes or less. Each one of these packets undergoes an encapsulation process. Under high network utilization or high network throughput, this may lead to high CPU utilization and high latency.

To reduce the CPU utilization and latency under such circumstances, the recommendations are

- Use Jumbo frames (i.e., MTU 9000 or more) to be able to send more data per packet and reduce the number of packets and overhead required to move the data from one Node to the other.
- Use NIC cards supporting VXLAN acceleration so the encapsulation process is offloaded to hardware and CPU.

**CAUTION** Not all the VXLAN-accelerated NICs support Jumbo frames. Consult the technical specification of your NIC provider. In those cases where VXLAN acceleration is not supported with Jumbo frames, the cluster administrators should avoid Jumbo frames as the NIC driver will determine the final behavior which may have a negative impact in performance.

## TUNING NETWORK DEVICES

Advanced Linux system administrator with deep understanding of the Linux networking stack and the available tuning options for high-performance



computing may use similar techniques with OpenShift clusters.

**CAUTION** Some of the following optimizations have limited availability or configuration options in some NIC drivers. Consult your hardware and driver technical information.

Some of the optimizations that may be considered are

- **Adjusting the number and size of RX and TX queues:** Improved throughput, latency, and multi-queue techniques can be used to distribute the processing of queues across multiple CPUs.
- **Interrupt coalescing:** Prevents interrupt storms and increases throughput or latency.
- **Adaptive RX and TX coalescing:** Interrupt delivery is optimized to improve latency or throughput based on packet rate.
- **Hardware-accelerated Receive Flow Steering (RFS) :** When supported by the NIC's driver, the NIC and the Kernel work together to determine which flows to send to which CPU for processing.
- **Adjusting IRQ affinity:** Optimizes for data locality for interrupts generated by the NIC.
- **Adjusting UDP receive queue size:** Increases throughput.
- **Generic Receive Offloading (GRO) and Large Receive Offloading (LRO)**
- **Receive Packet Steering (RPS) and Receive Side Scaling (RSS)**

Some of the Linux commands used for these optimizations are *sysctl* and *ethtool*.

These levels of optimizations are unique to each hardware and driver combination. As such, this book highlights the existence of these capabilities but leaves it to the reader to explore and test the ones suitable for their environment.



# Routing Optimizations

The *OpenShift Router* can handle the *Routes* for multiple applications. This can be anywhere from one to thousands of applications. The actual number of *Routes* an *OpenShift Router* can handle is determined by the technology in use by the applications behind the *Routes*.

As seen in previous chapters, the *OpenShift Router* is based on *HAProxy*. One of the tunable parameters for *HAProxy* is the *maxconn* parameter which is configurable by using the *ROUTER\_MAX\_CONNECTION*<sup>1</sup> environment variable of the *OpenShift Router DeploymentConfig*. This parameter sets the per-process maximum number of concurrent connections.

**NOTE** When configuring the *maxconn* parameter, consider *HAProxy* counts of the *frontend* connection and *backend* connection as two different connections. Because of this, a connection from an external client to an application load balanced by *HAProxy* counts as two.

Additional parameters<sup>2</sup> for the optimization of the *OpenShift Router* are

- CPU and interrupt affinity
- Increasing number of threads
- Setting up connection timeouts

**NOTE** *OpenShift Routers*, by default, listen on ports 80 (HTTP) and 443 (HTTPS), but they can be configured to listen for HTTP and HTTPS traffic on other ports. This option is configured using the environment variable *ROUTER\_SERVICE\_HTTP\_PORT* and the environment variable *ROUTER\_SERVICE\_HTTPS\_PORT*.

## ROUTE-SPECIFIC OPTIMIZATIONS ANNOTATIONS



In addition to the global configuration parameters of the HAProxy, OpenShift provides the ability to modify certain behavior on per-Route basis. This is done by using Route Annotations (see [Table 10-1](#)).

**Table 10-1** OpenShift Route Annotations <sup>3</sup>

Variable	Description
haproxy.router.openshift.io/balance	Load balancing algorithm: <i>source</i> , <i>roundrobin</i> , or <i>leastconn</i>
haproxy.router.openshift.io/disable_cookies	Disables the use of cookies to track relationships between connections
router.openshift.io/cookie_name	Optional cookie to use for Route

Variable	Description
haproxy.router.openshift.io/pod-concurrent-connections	Sets the maximum number of concurrent connections allowed for each back <i>Pod</i> from specific <i>Router</i>
haproxy.router.openshift.io/rate-limit-connections.concurrent-tcp	Limits the number of concurrent TCP connections by an IP address
haproxy.router.openshift.io/rate-limit-connections.rate-http	Limits the rate at which an IP address can make HTTP requests



Variable	Description
haproxy.router.openshift.io/rate-limit-connections.rate-tcp	Limits the rate at which an IP address can make TCP connections
haproxy.router.openshift.io/timeout	Sets a server-side timeout for the <i>Route</i>
router.openshift.io/haproxy.health.check.interval	Sets interval for the backend health checks
haproxy.router.openshift.io/ip_whitelist	(See “ <i>IP Whitelists</i> ” section)





Variable	Description
haproxy.router.openshift.io/hsts_header	Sets a Strict Transport Security header for the terminated or re-encrypted Route

# IP WHITELISTS

OpenShift supports the use of special annotations to restrict which source IP address or network can access a specific Route. The *ip\_whitelist* annotation (see Figure 10-2) is a space-separated list of whitelisted source IP addresses and CIDRs that are allowed to access the particular Route.

```
1 ...
2 metadata:
3   annotations:
4     haproxy.router.openshift.io/ip_whitelist: 209.132.183.105 192.168.1.0/24 10.5.25.0/24
5 ...
```

Figure 10-2 IP Whitelist annotations for a Route

# OPENSIFT ROUTER SHARDING

To horizontally scale the routing layer, OpenShift provides the capability to define and use Router Shards. In this case, the Routes are shared among a group of Routers based on a selection expression defining the Shard. There are two levels of Route sharding:



- Cluster administrators configure and manage sharding at cluster-wide level.
- Users can configure sharding for namespaces where they have admin privileges.

When using sharding, each *Router* in the group handles a portion of the traffic based on the assigned *Shard*.

**NOTE** Based on the *selection expression*, the *Router Shards* can be unique, in which case a *Route* belongs to only one *Shard*, or there can be overlapping in which case some *Routes* can belong to more than one *Shard*.

When using *Router Sharding*, the first *Route* matching a particular *Shard* reserves the right to exist on that *Shard* permanently and even across restarts. Figure 10-3 illustrates both ways of configuring *Router Sharding*.

USING NAMESPACE LABELS	USING ROUTE LABELS
<pre> 1 # Create new router (no replicas) 2 oc adm router router-dev --replicas=0 3 4 # Set the NAMESPACE selection expression 5 oc set env dc/router-dev NAMESPACE_LABELS="router=dev" 6 7 # Run one replica of the router 8 oc scale dc/router-ns-dev --replicas=1 9 10 # Label a NAMESPACE to be matched by selection expression 11 oc label namespace myproject "router=dev" </pre>	<pre> 1 # Create new router (no replicas) 2 oc adm router router-shard-2 --replicas=0 3 4 # Set the Shard selection expression 5 oc set env dc/router-shard-2 ROUTE_LABELS="bizunit=prod" 6 7 # Run one replica of the router 8 oc scale dc/router-shard-2 --replicas=3 9 10 # Label a ROUTE to be matched by selection expression 11 oc label route myroute "bizunit=prod" </pre>

**Figure 10-3** Router Sharding using Namespace or Route labels

**NOTE** When using the *Namespace* labels, the *Service Account* assigned to the *Router* must have *cluster-reader* permission to access the labels in the *Namespace*s.

## Supporting Non-HTTP/HTTPS/TLS Applications

There is a wide range of applications that cannot be classified as HTTP-, HTTPS-, or TLS-based applications. For example:

- Applications using specialized TCP protocols (i.e., database protocols)




- UDP-based applications
- Applications requiring direct access to the Pods IP

For these applications, OpenShift provides various mechanisms:

- Using IngressIP or ExternalIP
- Using NodePorts or HostPorts

## USING INGRESSIP AND EXTERNALIP

When using an ingressIP and externalIP, OpenShift uses Kube-Proxy to configure all Nodes into accepting traffic destined to the particular IP address. When traffic destined for a particular ExternalIP arrives to a Node, it forwards the traffic internally to the Pods associated to the Service (see #2 and #5 on Figure [10-4](#)).



**Figure 10-4** Traffic flow for non-http/https traffic

At first sight, both of these objects behave similarly but have a different default or intended purpose.

- **IngressIP:** This IP address is allocated from the *ingressIPNetworkCIDR* (default to 172.29.0.0/16 when not defined) for *Service* type *LoadBalancer*. This *CIDR* should not overlap with other IP ranges used in the Cluster.
- **ExternalIP:** This IP is allocated from a *CIDR* defined by the *externalIPNetworkCIDRs* variable in the `master-config.yaml` (see Figure 10-5). This can be a public IP address range or an organization-level visible and unique network *CIDR*.

**NOTE** IP Addresses from the *externalIP* *CIDR* are not managed by *OpenShift*. It is up to the network administrator to make sure the traffic destined to these IP arrives to the *Nodes*.

```
1 vi /etc/origin/master/master-config.yaml
2 ...
3 networkConfig:
4   externalIPNetworkCIDRs:
5     - 198.18.100.0/23
6     - <network_2>/<cidr>
7     - ...
8
9
10 # Restart Master service for changes to get into effect
11 $ master-restart api
12 $ master-restart controllers
```

**Figure 10-5** Defining an *externalIPNetworkCIDR*

Creating a service type *LoadBalancer* (see Listing 10-1) gets an *ingressIP* by default (see Figure 10-6).

```
apiVersion: v1
kind: Service
metadata:
  name: pgsql-lb
spec:
  ports:
    - name: pgsql
```



```
port: 5432
type: LoadBalancer
selector:
  name: postgresql
```

**Listing 10-1** Creating a Service type LoadBalancer

```
apiVersion: v1
kind: Service
metadata:
  name: postgresql-lb
spec:
  ports:
  - name: postgresql
    port: 5432
  type: LoadBalancer
  selector:
    name: postgresql
```

and

```
$ oc expose dc postgresql --type=LoadBalancer --name=postgresql-lb2
```

```
# psql -h 172.29.223.161 sampledadb pguser
Password for user pguser:
psql (9.2.24, server 9.6.10)
WARNING: psql version 9.2, server version 9.6.
Some psql features might not work.
Type "help" for help.

sampleddb=>
```

```
[root@ocp ~]# oc expose dc postgresql --type=LoadBalancer --name=postgresql-lb2
service/postgresql-lb2 exposed
[root@ocp ~]# oc get all
NAME                READY   STATUS    RESTARTS   AGE
pod/postgresql-1-vp6t 1/1     Running   0           4m

NAME                DESIRED   CURRENT   READY   AGE
replicationcontroller/postgresql-1 1 1 1 4m
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/postgresql	ClusterIP	172.30.240.202		5432/TCP	4m
service/postgresql-lb	LoadBalancer	172.30.106.105	172.29.121.182, 172.29.121.182	5432:30671/TCP	1m
service/postgresql-lb2	LoadBalancer	172.30.106.248	172.29.223.161, 172.29.223.161	5432:30318/TCP	4s

**Figure 10-6** LoadBalancer Service, IngressIP, and ExternalIP

The *LoadBalancer* resource can be created using a YAML file (see #1 in Figure 10-6) or using the OpenShift client command (see #2 in Figure 10-6). The resulting LoadBalancer object will be assigned an IP from the *ingressIPNetworkCIDR*.

**NOTE** The *CIDR* for *ingressIPNetworkCIDR* can be modified on the *master-config.yaml*.

Assigning an externalIP to a service is achieved by adding it to the *spec.externalIPs* definitions of the services (see Figure 10-7).

```
[root@ocp ~]# oc get svc
NAME                                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
glusterfs-dynamic-22a6a050-7efe-11e9-b815-001a0a160101 ClusterIP           172.30.55.81   <none>         1/TCP            32m
pgsql                               ClusterIP           172.30.240.202 <none>         5432/TCP         32m
pgsql-lb                            LoadBalancer       172.30.106.105 172.29.121.182, 172.29.121.182 5432:30671/TCP  30m
pgsql-lb2                           LoadBalancer       172.30.106.248 172.29.223.161, 172.29.223.161 5432:30318/TCP  28m
[root@ocp ~]# oc patch svc postgresql-lb2 -p '{"spec":{"externalIPs":["198.18.101.200"]}}'
service/postgresql-lb2 patched
[root@ocp ~]# oc get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
glusterfs-dynamic-22a6a050-7efe-11e9-b815-001a0a160101	ClusterIP	172.30.55.81	<none>	1/TCP	33m
pgsql	ClusterIP	172.30.240.202	<none>	5432/TCP	34m
pgsql-lb	LoadBalancer	172.30.106.105	172.29.121.182, 172.29.121.182	5432:30671/TCP	31m
pgsql-lb2	LoadBalancer	172.30.106.248	172.29.223.161, 198.18.101.200	5432:30318/TCP	30m

**Figure 10-7** Assigning externalIP to a Service



Another way to bring traffic into the Pods is by using a NodePort or HostPort. These two objects are similar in their behavior with respect to allocating ports in the actual Nodes. The difference is how the Ports are allocated in all Nodes from a range or allocated on the Node where the Pod is running.

- **NodePort:** Will allocate a port from the range 30000–32767 in all Nodes. (Note: It is possible to request a specific port in this range). The NodePort can be allocated for a Service or a specific Pod (see Figure 10-8).
- **HostPort:** Will allocate the specified port in the Node where it is running (see Figure 10-9).




Figure 10-8 Using NodePort

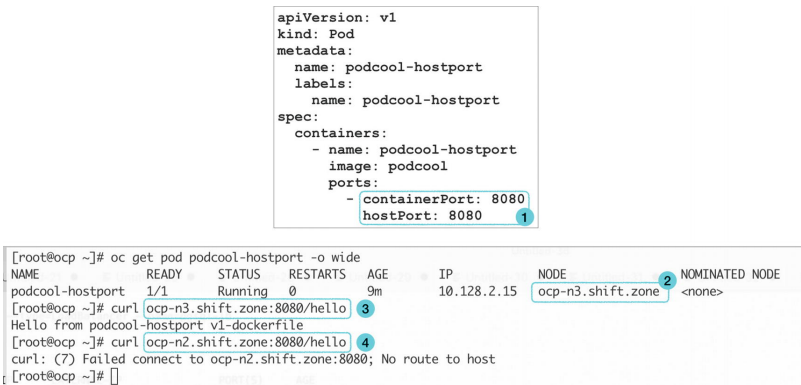


Figure 10-9 Using HostPort

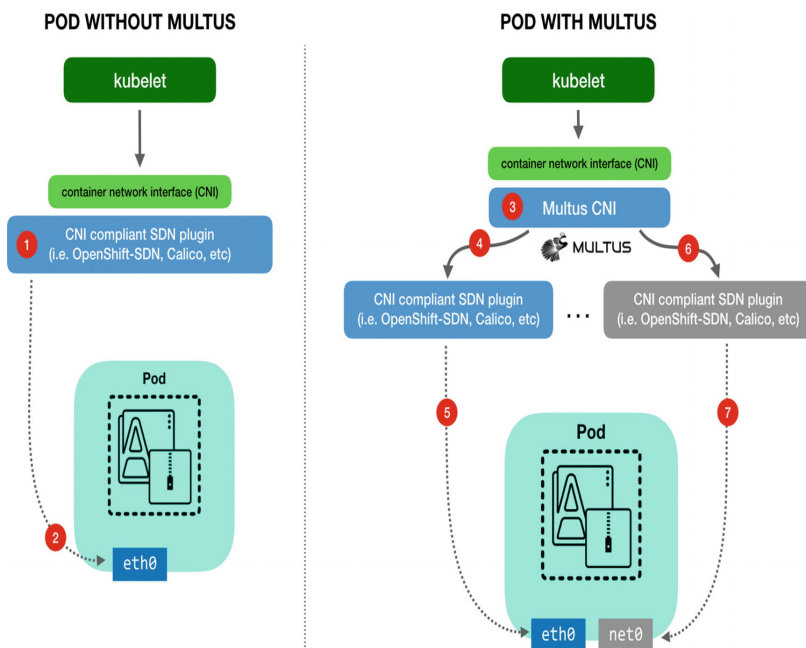
## Multiple NIC per POD

OpenShift 4.1 and later support the ability to provide multiple network interfaces to Pods. This capability is provided by the *Multus CNI*.



*Multus CNI* is a meta plugin for Kubernetes which enables the creation of multiple network interfaces per *Pod*. Each interface can be using a different CNI plugin.

As seen in Figure 10-9, when *Multus CNI* receives the request for the creation of a new network interface for the *Pod*, it sends that request to the primary Kubernetes CNI (see #4 in Figure 10-10) for the creation of the `eth0` interface. In addition, it interprets the *Pod annotations* to invoke additional *CNIs* to add other interfaces (see #6 in Figure 10-10) to the *Pod*.



**Figure 10-10** Multus CNI logical diagram

Multus requires the creation of a `NetworkAttachmentDefinition` defining the additional CNI (see #1 in Figure 10-11). The Pod must be annotated with the additional CNIs to use to provide additional interfaces (see #2 in Figure 10-11). At the Pod level, the new network interface is created (see #4 in Figure 10-11).



```

1  apiVersion: "k8s.cni.cncf.io/v1"
2  kind: NetworkAttachmentDefinition
3  metadata:
4    name: my-cni2-conf
5  spec:
6    config: '{
7      "cniVersion": "0.3.0",
8      "type": "macvlan",
9      "master": "eth0",
10     "mode": "bridge",
11     "ipam": {
12       "type": "host-local",
13       "subnet": "192.168.2.0/24",
14       "rangeStart": "192.168.2.10",
15       "rangeEnd": "192.168.2.200",
16       "routes": [
17         { "dst": "0.0.0.0/0" }
18       ],
19       "gateway": "192.168.2.1"
20     }
21   }'
22
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: podcool-multus
5    annotations:
6      k8s.v1.cni.cncf.io/networks: my-cni2-conf
7  spec:
8    containers:
9      - name: podcool-multus
10        image: podcool
11        ports:
12          - containerPort: 8080
13
1  $ oc exec -ti podcool-multus sh
2  /usr/src/app $ ip addr | grep -A2 "@
3
4  3: eth0@if12: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1450 qdisc noqueue state UP
5      link/ether 0a:58:0a:80:02:0f brd ff:ff:ff:ff:ff:ff
6      inet 10.128.2.15/23 brd 10.128.3.255 scope global eth0
7
8  4: net0@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UNKNOWN group default
9      link/ether 00:1a:4a:16:01:01 brd ff:ff:ff:ff:ff:ff link-netnsid 0
10     inet 192.168.2.17/24 scope global net0
11

```

**Figure 10-11** Defining NetworkAttachmentDefinition and using Multus CNI

## OpenShift ServiceMesh

The OpenShift ServiceMesh is based on the upstream project Maistra.<sup>5</sup>

Some of the components of OpenShift ServiceMesh are

- **Istio:** Based on the Istio<sup>6</sup> project; enables the intelligent control of the flow of traffic; enables the authentication, authorization, and encryption of communication between microservices; enforces policies; and enables observability of the communication among the microservices of an application
- **Envoy:** Service proxy used by Istio and based on Envoy Proxy<sup>7</sup> project
- **Jaeger:** Distributed tracing capability based on the Jaeger<sup>8</sup> project (see #2 in Figure 10-12)

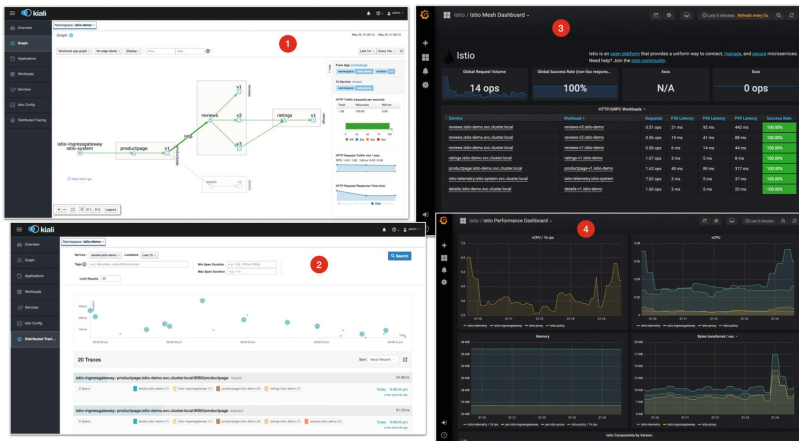




- **Kiali**<sup>9</sup>: Graphical interface integrating the components of OpenShift ServiceMesh (see #1 in Figure 10-12)
- **Grafana**: Used for the Istio mesh dashboards (see #3 and #4 in Figure 10-12)
- **Prometheus**: Used to collect Istio mesh metrics
- **Elasticsearch**: Used as the backend storage for the Istio metrics

**NOTE** At the time of this writing, installing OpenShift ServiceMesh<sup>10</sup> in OCP 3.11.x is still considered a *Technology Preview* capability.

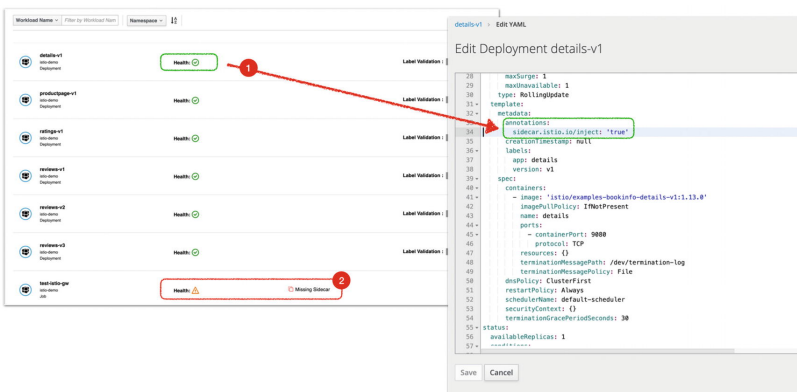
Some of the OpenShift ServiceMesh consoles are shown in Figure 10-12.



**Figure 10-12** OpenShift ServiceMesh

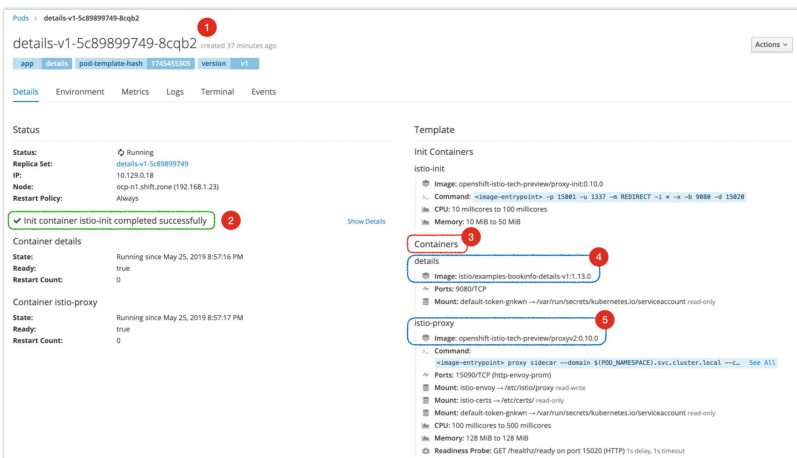
The main functionality of the OpenShift ServiceMesh requires injecting the Istio sidecar, the Envoy proxy, into the Pod. This requires the proper annotation of the *Deployment* configuration (see #1 in Figure 10-13).





**Figure 10-13** OpenShift ServiceMesh annotations for Istio sidecar

Once the Istio proxy sidecar is injected into the Pod (see #5 in Figure 10-14), all traffic incoming or outgoing to that Pod goes over the Istio-proxy sidecar container.



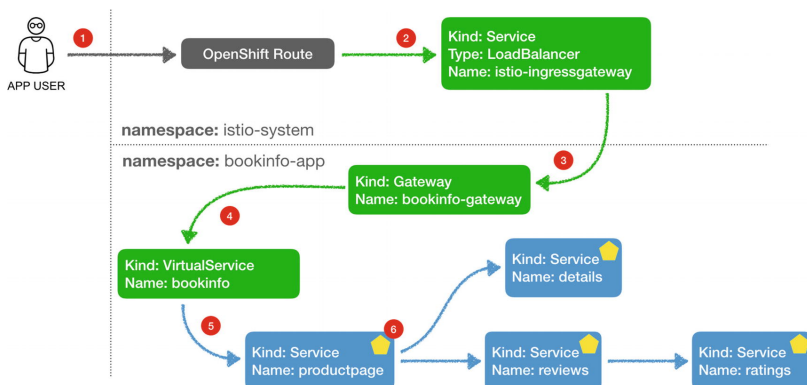
**Figure 10-14** OpenShift ServiceMesh Istio-Proxy sidecar injection

From the *OpenShift ServiceMesh* perspective, when using *OpenShift Routes* in conjunction with the *Istio Gateway* resources, the traffic flow will be as follows:

- External traffic arrives to the Route (see #1 in Figure 10-15) which points to a *LoadBalancer* type Service (see #2 in Figure 10-15).
- A *LoadBalancer* resource gets allocated an IngressIP and the cluster administrator could also assign an ExternalIP.



- Traffic is then delivered to the *Istio Gateway* on the destination *Project* or *Namespace* as seen in #3 on Figure [10-15](#).
- The *Istio Gateway* is considered the edge of the *Mesh* for incoming and outgoing connections. It describes the ports and protocol (HTTP/HTTPS/TCP) it will accept traffic for.
- Traffic accepted at the *Istio Gateway* is forwarded based on the *VirtualService* definition (see #5 on Figure [10-15](#)).
- The *VirtualService* defines one or more destinations where the traffic should go inside the *ServiceMesh* to reach the actual destination (i.e., a *Service* or *Pod*).
- Any *Service* or *Pod* annotated for the *Istio Proxy* will have the Istio sidecar injected into the Pods (see yellow pentagon shapes in Figure [10-15](#)).
- After this point, the metrics and visibility provided by the *Istio Proxy* are available over the *Kiali* console.



**Figure 10-15** Traffic flow with OpenShift ServiceMesh

**CAUTION** At the moment of this writing, the Istio-Proxy (Envoy) has limited support for non-TCP traffic. Applications relying on non-TCP protocols should investigate the impact of these limitations to avoid service disruption.

## Summary



This chapter covers some of the advanced network optimizations available in the OpenShift Container Platform (OCP). Some of these optimizations, like hardware acceleration, are dependent on the availability of underlying infrastructure supporting the capability. Other optimizations are more in the fine-tuning of configuration attributes to increase performance and scalability of the capability, like the optimizations available for the *OpenShift Routers*.

In addition to the optimization, this chapter describes the use of IngressIP, ExternalIPs, NodePorts, and HostPorts to bring specialized IP protocols to Services and Pods running on the platform.

Finally, the chapter explored advanced functionalities provided by OpenShift Multi-Network capabilities with Multus and the OpenShift ServiceMesh with Istio, Jaeger, and other upstream projects.

Some of the optimizations described in this chapter are intended for OpenShift bare-metal deployments. Chapter 11 provides a glimpse of the installation of OpenShift 4.1 using the User Provisioned Infrastructure (UPI) deployment option. This new OpenShift version provides the support for advanced networking capabilities like the multiple network for Containers using Multus, OpenShift ServiceMesh, and many others.

---

## Footnotes

- 1 Additional information about configuring the maximum number of connection is available at the online documentation:

[https://docs.openshift.com/container-platform/3.11/scaling\\_performance/routing\\_optimization.html#scaling-performance-optimizing-router-haproxy-maxconn](https://docs.openshift.com/container-platform/3.11/scaling_performance/routing_optimization.html#scaling-performance-optimizing-router-haproxy-maxconn)

- 2 OpenShift HAProxy optimization parameters:

[https://docs.openshift.com/container-platform/3.11/scaling\\_performance/routing\\_optimization.html#scaling-performance-optimizing-router-haproxy](https://docs.openshift.com/container-platform/3.11/scaling_performance/routing_optimization.html#scaling-performance-optimizing-router-haproxy)



3 Additional information and the updated list of possible Route

Annotations are available from the online documentation:

<https://docs.openshift.com/container-platform/3.11/architecture/networking/routes.html#route-specific-annotations>

4 Additional information on providing multiple network interface to Pods can be found at the OCP 4.1 online documentation:

<https://docs.openshift.com/container-platform/4.1/networking/managing-multinetworking.html>

5 For more details of the Maistra project, refer to the online documentation at <https://maistra.io/docs/>

6 Upstream Istio project is available at <https://istio.io>

7 Upstream Envoy Proxy project is available at [www.envoyproxy.io](http://www.envoyproxy.io) (<http://www.envoyproxy.io>)

8 Upstream Jaeger project is available at [www.jaegertracing.io](http://www.jaegertracing.io) (<http://www.jaegertracing.io>)

9 Upstream Kiali project is available at [www.kiali.io](http://www.kiali.io) (<http://www.kiali.io>)

10 The instructions for the installation and configuration of the OpenShift Service Mesh are available at the online documentation:

<https://docs.openshift.com/container-platform/3.11/servicemesh-install/servicemesh-install.html#installing-service-mesh>



 [PREV](#)  
[9. Day-2 Operations](#)

[NEXT](#)   
[11. OCP 4.1 UPI Mode Ba...](#)

