

INTRODUCTION TO THE VSPHERE CLUSTERING SERVICE (VCLS)

Table of Contents

[Overview](#)

[Basic Architecture](#)

[Cluster Service Health](#)

-

[Agent VM Resources](#)

[UI Overview](#)

[Operations](#)

- [Automation and vCLS](#)

Introduction to the vSphere Clustering Service (vCLS)

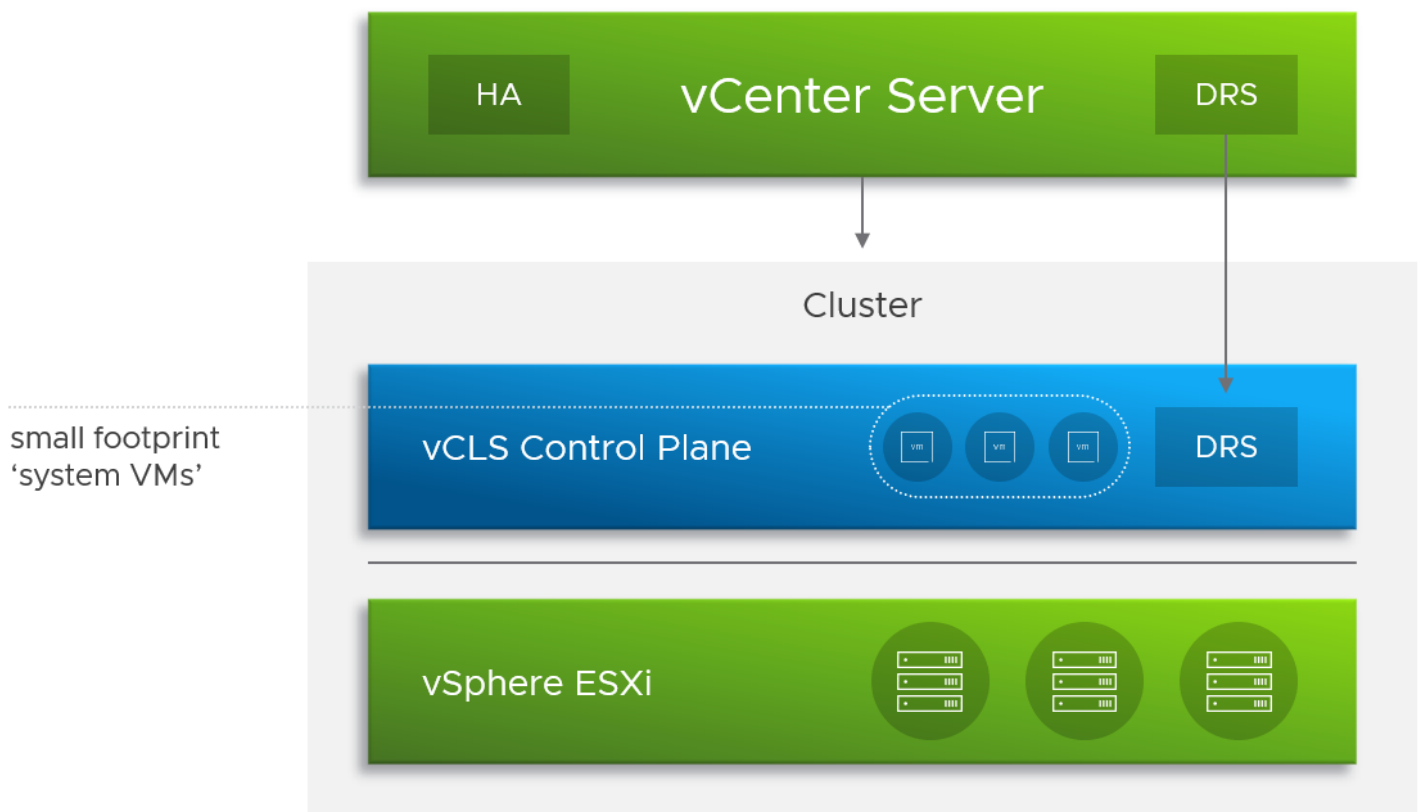
Overview

The vSphere Clustering Service (vCLS) is a new capability that is introduced in the vSphere 7 Update 1 release. It's first release provides the foundation to work towards creating a decoupled and distributed control plane for clustering services in vSphere.

The challenge being that cluster services, like the vSphere Distributed Resource Scheduler (DRS), depend on vCenter Server availability for its configuration and operation. And while there's ways to increase availability for vCenter Server, think about vSphere High Availability (HA) and vCenter Server High Availability (VCHA), its dependency is not ideal. Also, when thinking about vCenter Server scalability in large on-prem and public clouds, we need a better solution to support clustering services. That's why vCLS is introduced. In the first release, a subset of DRS capabilities is already using the new vCLS feature.


Basic Architecture

The basic architecture for the vCLS control plane consists of maximum 3 virtual machines (VM), also referred to as system or agent VMs which are placed on separate hosts in a cluster. These are lightweight agent VMs that form a cluster quorum. On smaller clusters with less than 3 hosts, the number of agent VMs is equal to the numbers of ESXi hosts. The agent VMs are managed by vSphere Cluster Services. Users are not expected to maintain the lifecycle or state for the agent VMs, they should not be treated like the typical workload VMs.



Cluster Service Health

The agent VMs that form the cluster quorum state, are self correcting. This means that when the agent VMs are not available, vCLS will try to instantiate or power-on the VMs automatically.



There are 3 health states for the cluster services:

- **Healthy** – The vCLS health is green when at least 1 agent VM is running in the cluster. To maintain agent VM availability, there's a cluster quorum of 3 agent VMs deployed.
- **Degraded** – This is a transient state when at least 1 of the agent VMs is not available but DRS has not skipped its logic due to the unavailability of agent VMs. The cluster could be in this state when either vCLS VMs are being re-deployed or getting powered-on after some impact to the running VMs.
- **Unhealthy** – A vCLS unhealthy state happens when a next run of the DRS logic (workload placement or balancing operation) skips due to the vCLS control-plane not being available (at least 1 agent VM).

Agent VM Resources

The vCLS agent VMs are lightweight, meaning that resource consumption is kept to a minimum. vCLS automatically creates a max of 3 agent VMs per cluster in an existing deployment when vCenter Server is upgraded to vSphere 7 update 1. In a greenfield scenario, they are created when ESXi hosts are added to a new cluster. If no shared storage is available, the agent VMs are placed on local storage. If a cluster is formed before shared storage is configured on the ESXi hosts, as would be the case when using vSAN, it is strongly recommended to move the vCLS agent VMs to shared storage after.

The agent VMs run a customized Photon OS. The resource specification per agent VM is listed in the following table:

Memory	128 MB
Memory reservation	100 MB
Swap Size	256 MB
vCPU	1
vCPU reservation	100 MHz
Disk	2 GB
Ethernet Adapter	-
Guest VMDK Size	~245 MB
Storage Space	~480 MB

The 2 GB virtual disk is thin provisioned. Also, there's no networking involved, so there's no network adapter configured. The agent VMs are not shown in the *Hosts and Clusters* overview in the vSphere Client. The *VMs and Templates* view now contains a new folder, vCLS, that contains all vCLS agent VMs. With multiple clusters, all vCLS agent VMs will show, numbered consecutively.

UI Overview

The screenshot displays the vSphere Client interface for a vCLS (1) VM. The left-hand navigation pane shows a tree structure under '10.196.180.159' > 'datacenter' > 'Discovered virtual machine' > 'vCLS', with 'vCLS (1)', 'vCLS (2)', and 'vCLS (3)' listed. The main content area is titled 'vCLS (1)' and includes tabs for Summary, Monitor, Configure, Permissions, Datastores, Networks, Snapshots, and Updates. The 'Summary' tab is active, showing the VM is 'Powered On'. Key details include: Guest OS: Other 4.x or later Linux (64-bit); Compatibility: ESXi 6.0 and later (VM version 11); VMware Tools: Running, version:11296 (Guest Managed); DNS Name: photon3-hdcs; IP Addresses: 10.196.180.157; Host: 10.196.180.157; Managed By: VMware vCenter Server. Resource usage on the right shows CPU Usage at 0 Hz, Memory Usage at 0 B, and Storage Usage at 1,008 MB. A warning banner states: 'vSphere Cluster Service VM is required to maintain the health of vSphere Cluster Services. Power state and resource of this VM is managed by vSphere Cluster Services.' The 'VM Hardware' section lists 1 CPU(s), 0.13 GB memory (0 GB active), and a 2 GB hard disk. The 'Notes' section contains a detailed message about the VM's role in maintaining cluster health.

The vSphere Client contains messages and notes to show information about the vCLS agent VMs, also stating that the power state and resources of these VMs is handled by the vCLS.

Operations

As stated before, the agent VMs are maintained by vCLS. There's no need for VI admins to power-off the VMs. In fact, the vSphere Client shows a warning when a agent VM is powered-off.

Confirm Power Off | vCLS (3)



vSphere Cluster Services manage the resources, power state and availability of this VM. vSphere Cluster Service VMs are required for maintaining the health and availability of vSphere Cluster Services. Any impact on the power state or resources of these VMs might degrade the health of vSphere Cluster Services and cause vSphere DRS to cease operation for the cluster.


For more information see KB 79892.

Power off the selected virtual machine?


NO

YES

When a host is placed into maintenance mode, the vCLS agent VMs are migrated to other hosts within the cluster like regular VMs. Customers should refrain from removing, or renaming the agent VMs or its folder in order to keep the cluster services healthy.


vSphere ESX Agent Manager
ACTIONS ▾

Summary
Monitor
Configure
VMs



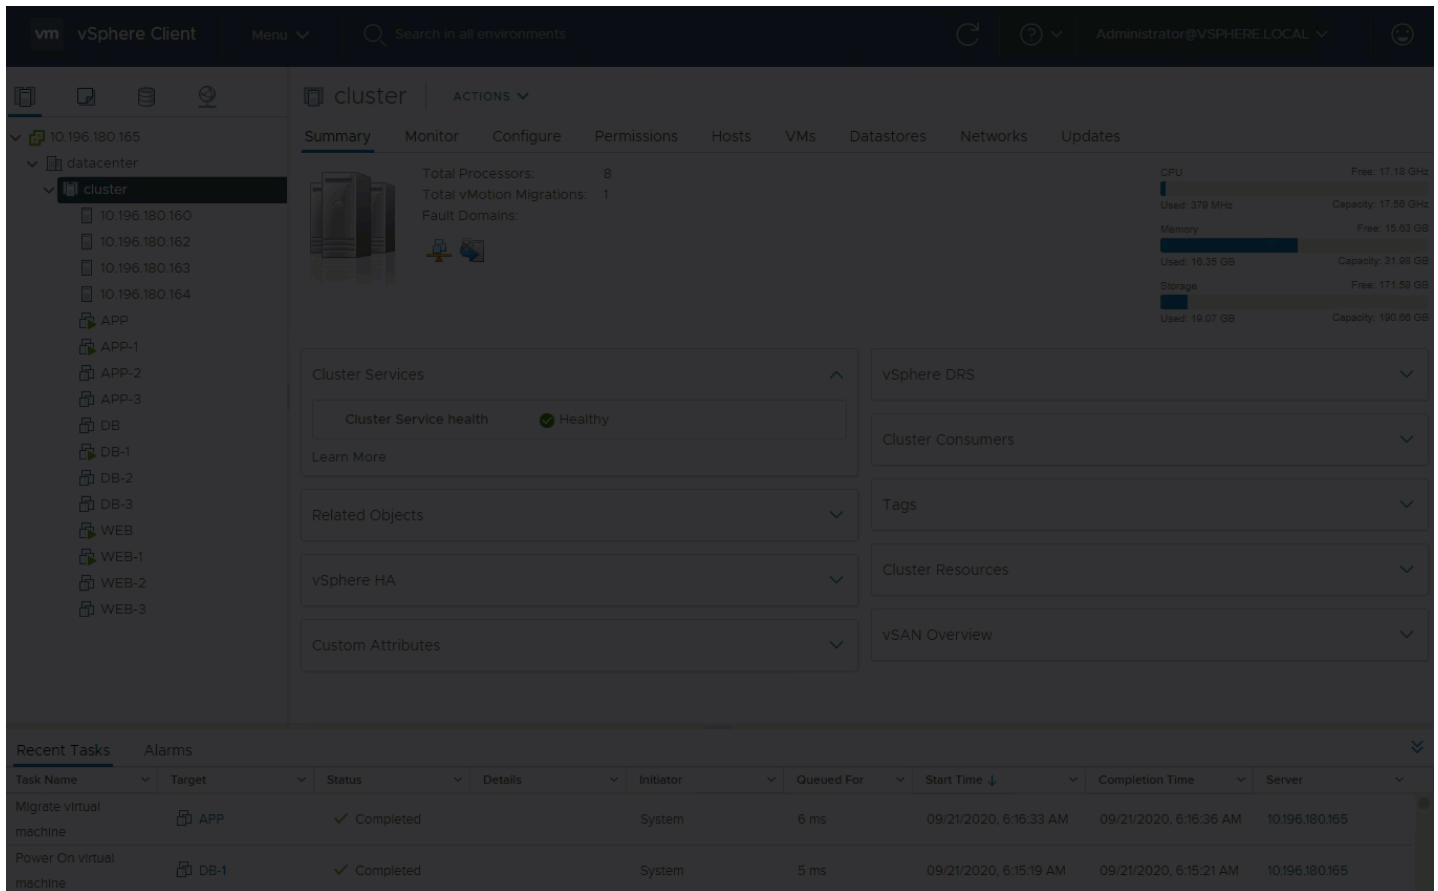
Product: vSphere ESX Agent Manager
Vendor: VMware Inc.
Version: 7.0.1
vCenter Server: 10.196.180.159

Details
⤴

Description	Manages the life-cycle of agent VIBs and VMs
Product Page	http://www.vmware.com/
Vendor Page	http://www.vmware.com/

The lifecycle for vCLS agent VMs is maintained by the vSphere ESX Agent Manager (EAM). The Agent Manager creates the VMs

automatically, or re-creates/powers-on the VMs when users try to power-off or delete the VMs. In the example below, you'll see a power-off and a delete operation. Both from which the EAM recovers the agent VM automatically.



Automation and vCLS

For customer using scripts to automate tasks, it's important to build in awareness to ignore the agent VMs in, for example clean-up scripts to delete stale VMs. Identifying the vCLS agent VMs is quickly done in the vSphere Client where the agent VMs are listed in the vCLS folder. Also, examining the VMs tab under Administration > vCenter Server Extensions > vSphere ESX Agent Manager lists the agent VMs from all clusters managed by that vCenter Server instance.

Every agent VM has additional properties so they can be ignored with specific automated tasks. These properties can also be found using the Managed Object Browser (MOB). The specific properties include:

- ManagedByInfo
 - extensionKey == "com.vmware.vim.eam"
 - type == "cluster-agent"
- ExtraConfig keys
 - "eam.agent.ovfPackageUrl"
 - "eam.agent.agencyMold"
 - "eam.agent.agentMold"

Home		
Data Object Type: ManagedByInfo		
Parent Managed Object ID: vm-1004		
Property Path: config.managedBy		
Properties		
NAME	TYPE	VALUE
extensionKey	string	"com.vmware.vim.eam"
type	string	"cluster-agent"

Home		
Data Object Type: OptionValue		
Parent Managed Object ID: vm-1004		
Property Path: config.extraConfig["HDCS.agent"]		
Properties		
NAME	TYPE	VALUE
key	string	"HDCS.agent"
value	string	"true"

vCLS Agent VMs have an additional data property key "HDCS.agent" set to "true". This property is automatically pushed down to the ESXi host along with the other VM ExtraConfig properties explicitly set by EAM.



**VMware, Inc. 3401 Hillview Avenue Palo Alto CA 94304 USA Tel 877-486-9273 Fax
650-427-5001 www.vmware.com**

Copyright © 2019 VMware, Inc. All rights reserved. This product is protected by U.S. and international copyright and intellectual property laws. VMware products are covered by one or more patents listed at <http://www.vmware.com/go/patents>. VMware is a registered trademark or trademark of VMware, Inc. in the United States and/or other jurisdictions. All other marks and names mentioned herein may be trademarks of their respective companies.